

Network-Validated Distributed ADMM Coordination of Cold-Climate Electric Heating under Communication Failures

NILSON HENAO

University of Québec at Trois-Rivières, Trois-Rivières, QC G8Z 4M3, Canada (e-mail: nilsonfh@gmail.com)

NOMENCLATURE

N	Number of heating agents (homes), $N = 30$.
T	Number of time steps over the day, $T = 96$.
$h_{i,t}$	Baseline (thermostat-driven) heating power of home i at step t .
$x_{i,t}$	Coordinated heating power of home i at step t .
$\hat{x}_{i,t}$	Imputed heating power of a non-responsive home i .
B_t	Fixed (non-deferrable) background aggregate load at step t .
c	Constant flat-aggregate target level (energy-conserving).
α	Per-step comfort deferrability band, $\alpha = 0.5$.
λ	Comfort-deviation weight, $\lambda = 0.15$.
μ	Aggregate-flattening weight, $\mu = 1.0$.
ρ_{ADMM}	ADMM penalty parameter, $\rho_{\text{ADMM}} = 1.0$.
ϱ	Non-responsive (communication-failure) fraction of agents.
z	ADMM auxiliary shared-aggregate variable.
u	ADMM scaled dual variable.
ν_i	Scalar dual of the daily-energy hyperplane in the projection.
$V_{\min}$	Worst-of-day minimum bus voltage (p.u.).

I. INTRODUCTION

THE electrification of heating in cold-climate regions concentrates distribution-system stress into a few winter hours. In the province of Québec, where electric space heating is the dominant heating technology, the system peak is driven almost entirely by simultaneous resistive and heat-pump heating on the coldest mornings and evenings. As distributed energy resources (DERs) proliferate and heating loads grow, these synchronized peaks threaten thermal limits on feeder trunk conductors and depress voltages far from the

substation. Because the flexibility lives behind thousands of meters and not at a single controllable node, centralized dispatch scales poorly; decentralized coordination, in which each home solves a small local problem and exchanges only an aggregate signal, is the natural fit.

The alternating direction method of multipliers (ADMM) [1] is a workhorse for this setting: its sharing form lets a coordinator flatten a common resource (here, total feeder load) while each agent retains private control of its own heating schedule. ADMM-based demand response is well studied, yet two gaps undermine its real-world credibility.

The first gap is fragility to communication failures. Multi-agent coordination implicitly assumes every agent reports on every iteration; when a fraction of agents fall silent, the aggregate the coordinator optimizes is no longer the aggregate that materializes on the wire. How badly the coordinated outcome degrades — and how an operator might recover it by forecasting the missing contributions — is rarely quantified end to end.

The second gap is physical validity. The literature typically stops at the coordinated aggregate kilowatt signal and reports peak or peak-to-average reductions. It almost never asks the operational question that matters: is the coordinated aggregate actually feasible on the feeder? A flatter aggregate is worthless if it still overloads a trunk line or sags a remote bus below limits.

This paper closes the loop between coordination and feeder physics. We adapt the sharing-ADMM-plus-imputation methodology from a prior multi-agent thesis and extend it with an AC power-flow validation stage, turning a purely signal-level study into a network-validated one. Our contributions are:

- A sharing-ADMM coordinator for deferrable electric heating that flattens the feeder aggregate subject to per-

home comfort bounds and a daily-energy conservation constraint (load is shifted, never curtailed).

- A LightGBM imputer that reconstructs the heating of non-responsive agents from ambient temperature, time-of-day, and nameplate features, letting the coordinator stay accurate as the non-responsive fraction grows.
- An AC power-flow validation stage on the IEEE 33-bus feeder in pandapower that distributes the coordinated aggregate across load buses, calibrates trunk ampacity to a realistic uncoordinated overload, and records worst-of-day minimum voltage and maximum line loading — closing the loop on physical feasibility.
- A fully reproducible, governed study driving the entire pipeline from a declarative configuration, with every reported number traceable to a recorded artifact.

The central finding is that coordination not only flattens the aggregate (a 12.7% peak cut) but clears a 115% thermal overload to 99.7%, and that this cleared state holds under communication failure up to a non-responsive fraction of 0.6 before re-violating at 0.8. Throughout, we keep the communication-failure fraction (denoted ϱ) notationally distinct from the ADMM penalty (ρ_{ADMM}) so the two roles of the symbol are never conflated; the Nomenclature collects all symbols. We also stress two honest caveats that the results section quantifies: the flattening preserves per-home comfort by construction (heat is re-timed, never withheld), and because the objective targets a flat profile rather than the time-of-use price, the energy-bill saving is deliberately modest — the value of this demand response is capacity relief, demonstrated here as restored feeder feasibility, not a cheaper bill.

The remainder of the paper is organized as follows. Section IV is preceded by the related work and the system model; Section II surveys ADMM-based demand response, multi-agent resilience, telemetry imputation, and distribution power flow, and states our delta against each. Section III formalizes the agents, the feeder, and the communication-failure model. Section IV develops the sharing-ADMM updates, the LightGBM imputer, and the power-flow validation loop. Section V details the case-study configuration, Section VI reports the aggregate, robustness, and network-feasibility results, and Section VII concludes.

II. RELATED WORK

Our contribution sits at the intersection of five strands of literature: distributed (ADMM/consensus) coordination, network-aware demand response, resilience of distributed coordination to communication failures, machine-learning imputation of missing load telemetry, and the flexibility of residential electric heating. We review each and then make the gap explicit (Table 1).

A. ADMM AND CONSENSUS FOR DISTRIBUTED ENERGY COORDINATION

ADMM [1] decomposes a coupled optimization into local subproblems linked by a consensus or sharing constraint,

and is now a standard tool for distributed energy coordination. The sharing form (Boyd et al., Section 7.3) is particularly suited to demand flattening, where many agents contribute to one shared resource. It underpins distributed optimal power flow [2], communication-efficient and consensus-based demand response [3], [4], asynchronous microgrid energy management [5], and hierarchical transactive coordination of home energy management systems [6]. These works establish that ADMM converges to a flattened or cost-optimal aggregate; with few exceptions they treat the aggregate signal as the deliverable and do not verify that it is feasible on the underlying feeder.

B. NETWORK-AWARE COORDINATION

A second strand embeds network physics in the coordination itself. Distributed AC optimal power flow couples ADMM with the branch-flow equations [7], and recent work coordinates distributed energy resources with inverter Volt-Var control through a three-block AC-OPF ADMM [8]. Network-constrained demand-response pricing [9], ADMM-based congestion-management markets [10], and capacity-constrained electric-vehicle charging [11] all respect feeder limits during coordination. A parallel line handles uncertainty through chance-constrained AC-OPF, enforcing probabilistic voltage and thermal limits for renewables [12] and flexible-load ensembles [13]. This literature proves that network-feasible coordination is achievable, but each work assumes perfect, complete communication: every agent reports on every iteration.

C. RESILIENCE TO COMMUNICATION FAILURES

A third strand studies how coordination degrades when agents fall silent. Brown et al. show that even a single lost link can drive multi-agent systems to arbitrarily poor states [14], and fault-tolerant consensus is a mature survey topic [15]. In distributed optimization, stochastic communication delay slows ADMM-based OPF [16], lossy links are handled by robust dispatch [17] and relaxed asynchronous ADMM with convergence guarantees [18], and adversarial or non-responsive agents can be detected and isolated to preserve convergence [19]. Learning has begun to enter the loop — e.g. to warm-start consensus-ADMM OPF [20]. This literature characterizes when coordination breaks and offers robust-by-design or detect-and-isolate responses, but it rarely reconstructs a silent agent's contribution, and never propagates that reconstruction error into feeder physics.

D. IMPUTATION AND FORECASTING OF MISSING TELEMETRY

When an agent falls silent the coordinator must estimate its consumption. Machine-learning imputation of missing smart-meter data spans denoising diffusion models [21], correlation-clustering for faulted grid data [22], and low-rank recovery of asynchronous meter streams [23]. For the closely related task of short-term load prediction, gradient-boosted trees are competition-winning regressors [24], [25],

motivating our use of LightGBM [26] to impute a silent home's heating from weather and calendar features. Crucially, in all of this work the imputer or forecaster is evaluated in isolation — by held-out RMSE — and is never embedded in a distributed coordinator nor propagated into a power-flow feasibility check.

E. RESIDENTIAL ELECTRIC-HEATING FLEXIBILITY

Finally, the resource we coordinate — thermostatically controlled electric heating — is itself well studied. Aggregate thermostatic flexibility can be quantified for demand response [27], direct electric space heating combined with partial thermal storage shifts load under comfort constraints via linear programming [28], heat-pump-plus-storage flexibility reduces peak demand and generation capacity [29], and field experiments confirm the aggregate effect and its rebound dynamics [30]. In cold climates specifically, electrically heated buildings with thermal storage enable peak shaving [31], and multi-unit residential flexibility and predictive heating control have been demonstrated in Québec-style winter-peaking contexts [32], [33]. The distributed-control variants closest to ours coordinate large TCL populations through asynchronous greedy [34] or resilient-consensus [35] schemes — yet these remain abstract control formulations without an AC feeder model, and the broader thermal-DR literature is overwhelmingly centralized and not network-validated.

F. THE GAP

Table 1 positions representative prior work along the five capabilities our study requires together: distributed ADMM/consensus coordination, AC power-flow validation, an explicit communication-failure model, machine-learning imputation of the missing agents, and uncertainty quantification of the resulting risk. Every ingredient exists — but only in disjoint subsets. Distributed-ADMM coordinators assume perfect communication; network-aware and chance-constrained methods are centralized or assume complete reporting; resilience methods detect or tolerate failures rather than reconstruct the missing agents; imputation accuracy is studied standalone; and thermal-DR studies seldom touch the feeder at all. To our knowledge no prior work closes the full loop — distributed ADMM coordination of cold-climate electric heating, machine-learning imputation of communication-failed homes, AC power-flow validation on a distribution feeder, and Monte-Carlo quantification of the resulting violation probability under a swept non-responsive fraction. That closed loop is the contribution of this paper.

III. SYSTEM MODEL

A. AGENTS AND DEFERRABLE HEATING

We model $N = 30$ representative cold-climate electric-heating homes over a single synthetic Trois-Rivières cold day at 15-minute resolution ($T = 96$ time steps). Each home i has a thermostat-driven baseline heating profile $h_{i,t}$ produced by a behavioral model conditioned on the ambient

temperature, plus a fixed (non-deferrable) background load. The coldest-day minimum temperature is -21.9°C (daily mean -18.3°C), which sets a demanding heating duty.

Heating is treated as deferrable within a comfort band of $\pm 50\%$ of the baseline ($\alpha = 0.5$): the coordinator may shift when a home heats but must conserve the home's total daily heating energy, so demand is shifted rather than curtailed. With $x_{i,t}$ the coordinated heating power of home i at step t , each agent is constrained by

$$(1 - \alpha) h_{i,t} \leq x_{i,t} \leq (1 + \alpha) h_{i,t}, \quad \sum_t x_{i,t} = \sum_t h_{i,t}. \quad (1)$$

The box bound preserves thermal comfort at every step, and the equality plane preserves daily heating energy. Background load is held fixed and is not subject to (1).

B. THE FEEDER

Physical feasibility is assessed on the IEEE 33-bus radial test feeder [36] as implemented in pandapower [37] (the `case33bw` case): 33 buses, 32 lines, 32 distributed loads, a single external-grid substation at the head, and a nominal voltage of 12.66 kV. The coordinated aggregate is allocated to the 32 load buses in proportion to their native load shares and scaled so that the uncoordinated winter peak reaches the feeder. The trunk conductor ampacity is calibrated so that the uncoordinated peak loads the worst line to 115% — a realistic overload that motivates demand response in the first place. All loads use a power factor of 0.95.

C. COMMUNICATION-FAILURE MODEL

At coordination time a fraction ϱ of the agents become non-responsive and report nothing to the coordinator. We use ϱ for this communication-failure fraction throughout, kept notationally distinct from the ADMM penalty ρ_{ADMM} of Section IV. We sweep $\varrho \in \{0.0, 0.2, 0.4, 0.6, 0.8\}$, corresponding to 0, 6, 12, 18, and 24 silent homes out of 30. For each non-responsive agent the coordinator substitutes an imputed heating profile $\hat{x}_{i,t}$ produced by a trained regressor (Section IV) and proceeds as if it were the true contribution. The imputation error therefore propagates first into the coordinated aggregate and ultimately into the feeder power flow, which is exactly the propagation path this study quantifies.

IV. METHODOLOGY

A. SHARING-ADMM COORDINATION

Let $x_i \in \mathbb{R}^T$ be the heating schedule of home i , h_i its baseline, and $B \in \mathbb{R}^T$ the fixed background aggregate. The coordinator flattens the total feeder load toward a constant target c while respecting each home's comfort. We solve the sharing problem

$$\min_{\{x_i\}} \sum_{i=1}^N \frac{\lambda}{2} \|x_i - h_i\|_2^2 + \frac{\mu}{2} \left\| \sum_{i=1}^N x_i + B - c \mathbf{1} \right\|_2^2 \quad (2)$$

TABLE 1. Positioning against representative prior work. ✓ : addressed; ~: partial; –: not addressed. PF: power flow.

Representative work	Distributed ADMM/consensus	AC PF validation	Comms-failure model	ML imputation of missing agents	Uncertainty quantification
Distributed AC-OPF [2], [7]	✓	✓	–	–	–
ADMM AC-OPF DER coord. [8]	✓	✓	–	–	–
Consensus / sharing DR [4], [6]	✓	–	–	–	–
Chance-constrained OPF [12], [13]	–	✓	–	–	✓
ADMM-OPF under delay [16]	✓	✓	✓	–	~
Resilient distributed opt. [19]	✓	–	✓	–	–
Learning-augmented ADMM-OPF [20]	✓	✓	–	~	–
Distributed TCL control [34], [35]	✓	–	~	–	–
ML telemetry imputation [21], [23]	–	–	–	✓	–
Thermal-DR flexibility [27], [28]	–	–	–	–	–
This work	✓	✓	✓	✓	✓

Algorithm 1: Sharing-ADMM heating coordination

Input: baselines $\{h_i\}$, background B , target c , weights λ, μ , penalty ρ_{ADMM}
Output: coordinated schedules $\{x_i\}$
initialize $x_i \leftarrow h_i, z \leftarrow 0, u \leftarrow 0$;
repeat
 foreach home i (in parallel) **do**
 proximal step on $\frac{\lambda}{2}\|x_i - h_i\|^2$ plus penalty toward $z - u$;
 project x_i onto the box-and-energy set of (1);
 $z \leftarrow$ closed-form average enforcing the flat aggregate target;
 $u \leftarrow u + (\sum_i x_i + B) - z$;
until aggregate residual converged;
return $\{x_i\}$;

subject to (1) for every i , where c is the constant level that conserves total daily energy, $\lambda = 0.15$ weights comfort deviation, $\mu = 1.0$ weights flattening, and $\mathbf{1}$ is the all-ones vector. The first term keeps each home close to its comfortable baseline; the second penalizes any departure of the total load from a flat profile.

Following the sharing form of ADMM [1], we introduce an auxiliary aggregate variable z and a scaled dual u with penalty $\rho_{\text{ADMM}} = 1.0$, and iterate three steps. The **x -update** solves each home’s local problem in parallel: a comfort-plus-penalty proximal step followed by projection of the result onto the box-and-energy-plane feasible set of (1). The **z -update** is the closed-form average that drives the shared aggregate toward the flat target. The **u -update** is the standard scaled dual ascent on the sharing residual. Algorithm 1 summarizes one full sweep.

The x -update projection is what makes the scheme physically meaningful: it guarantees that every iterate respects comfort and conserves daily energy, so the coordinated schedule is feasible for the homes by construction. The number of iterations to convergence is recorded as a tractability metric.

1) Capped-energy projection in the x -update

We make the local step explicit because its structure is what keeps every iterate feasible. For home i the unconstrained comfort-plus-penalty objective is a separable strictly convex quadratic, so its minimizer is available in closed form. Writing $a_i = z - u$ for the current penalty anchor, the prox optimum is the per-step convex combination

$$\tilde{x}_i = \frac{\lambda h_i + \rho_{\text{ADMM}} a_i}{\lambda + \rho_{\text{ADMM}}}, \quad (3)$$

which trades the comfort pull toward h_i against the consensus pull toward a_i . This unconstrained point need not respect (1), so we Euclidean-project it onto the intersection of the comfort box $[(1 - \alpha)h_{i,t}, (1 + \alpha)h_{i,t}]$ and the daily-energy hyperplane $\sum_t x_{i,t} = \sum_t h_{i,t}$:

$$x_i = \arg \min_x \frac{1}{2}\|x - \tilde{x}_i\|_2^2 \quad \text{s.t.} \quad x \in \text{box}_i, \quad \mathbf{1}^\top x = \mathbf{1}^\top h_i. \quad (4)$$

Introducing a single scalar dual ν_i for the energy constraint, the projection has the water-filling form $x_{i,t} = \text{clip}_{[(1-\alpha)h_{i,t}, (1+\alpha)h_{i,t}]}(\tilde{x}_{i,t} + \nu_i)$, and the value of ν_i that restores $\mathbf{1}^\top x_i = \mathbf{1}^\top h_i$ is found by a 1-D bisection on the scalar dual. Because the clipped sum is monotone non-decreasing in ν_i , the bisection is exact and costs only a handful of evaluations, so the whole x -update is closed-form up to a scalar root-find. The subsequent z -update is fully separable per time step — each component of z is the average of the corresponding aggregate-plus-dual component driven toward the flat target c — and therefore also closed form, which is why a full ADMM sweep is inexpensive even with $N = 30$ agents over $T = 96$ steps.

B. IMPUTATION OF NON-RESPONSIVE AGENTS

When agent i is non-responsive its baseline h_i is unavailable to the coordinator, so we replace it with a learned estimate $\hat{h}_i$. We train a single LightGBM gradient-boosted-tree regressor [26] that maps a small feature vector — ambient temperature, the sine and cosine of the hour of day, and the agent’s nameplate heating level — to instantaneous heating power. The model is trained on the responsive population and applied step-by-step to each silent agent, yielding $\hat{x}_{i,t}$

that the coordinator treats as the agent's contribution. Generalization to unseen homes is assessed by a 5-fold leave-agents-out cross-validation, which holds out entire agents rather than random samples and therefore measures the realistic case of imputing a home the model never saw.

C. POWER-FLOW VALIDATION LOOP

The novel stage closes the loop on physics. For each scenario (uncoordinated, ideal, and each ϱ) we take the resulting aggregate, distribute it across the 32 IEEE-33 load buses by their native load shares, and scale so that the uncoordinated peak reaches the calibrated feeder. The trunk ampacity is fixed so the uncoordinated peak loads the worst line to 115%. We then run an AC Newton-Raphson power flow at every one of the 96 time steps (672 solves across all scenarios, all converged) and record, per scenario, the worst-of-day minimum bus voltage and the worst-of-day maximum line loading. A scenario is flagged as violating if any line exceeds 100% loading at any step. This turns an abstract kilowatt signal into a binary feeder verdict and a pair of continuous severity measures.

Three modeling choices deserve comment. Load allocation by native share keeps the spatial signature of the IEEE-33 case intact, so the worst line is the one the benchmark already stresses (the trunk leaving the substation) rather than an artifact of our placement. Ampacity calibration to 115% is what makes the experiment decision-relevant: a feeder that is comfortably within limits in the uncoordinated case has nothing to gain from demand response, so we set the binding constraint exactly where coordination must act to clear it. Solving every step rather than only the aggregate-peak step matters because imputation error can shift the worst instant: the uncoordinated worst is the evening peak, but under heavy imputation the binding step can migrate, and an aggregate-peak-only check would miss it. Reporting the worst-of-day envelope across all 96 steps therefore gives a conservative, instant-agnostic verdict.

V. CASE STUDY SETUP

The study runs on a single synthetic Trois-Rivières cold day with a minimum ambient temperature of -21.9°C , discretized at 15-minute resolution into $T = 96$ steps. We coordinate $N = 30$ electric-heating homes with comfort band $\alpha = 0.5$, comfort weight $\lambda = 0.15$, flattening weight $\mu = 1.0$, and ADMM penalty $\rho_{\text{ADMM}} = 1.0$. The communication-failure fraction is swept over $\varrho \in \{0.0, 0.2, 0.4, 0.6, 0.8\}$.

Energy cost is evaluated against a time-of-use tariff with an off-peak rate of 0.078, a mid-peak rate of 0.128, and an evening-peak rate of 0.158 CAD/kWh. Physical feasibility uses the IEEE 33-bus feeder with all loads at a power factor of 0.95, with the aggregate scaled so the uncoordinated winter peak reaches the feeder and trunk ampacity calibrated to a 115% uncoordinated overload. The software stack is pandapower 3.4.0 for power flow, scikit-learn 1.6.1 for the cross-validation harness, and LightGBM 4.6.0 for the

TABLE 2. Principal case-study parameters.

Parameter	Value
Homes N	30
Horizon / resolution	96 steps / 15 min
Coldest-day min. temperature	-21.9°C
Comfort band α	0.5
Comfort weight λ	0.15
Flattening weight μ	1.0
ADMM penalty ρ_{ADMM}	1.0
Non-responsive fraction ϱ	$\{0.0, 0.2, 0.4, 0.6, 0.8\}$
TOU off / mid / peak (CAD/kWh)	0.078 / 0.128 / 0.158
Feeder	IEEE 33-bus, 12.66 kV
Power factor	0.95
Uncoordinated worst-line loading	115% (calibrated)
Power flow	AC Newton-Raphson
Software	pandapower 3.4.0, scikit-learn 1.6.1, LightGBM

imputer. Every input, intermediate, and reported figure is emitted as a governed artifact so the run is reproducible from its declarative configuration. Table 2 collects the principal parameters.

A. REPRODUCIBILITY AND GOVERNED ARTIFACTS

The study is fully declarative. A `project.yaml` (a StudyProject contract) and a `workflow.yaml` (an eight-stage directed acyclic graph) drive the entire pipeline; no step depends on hand-edited intermediate state. The runner executes the stages in topological order — synthetic cold-day weather and per-home baseline generation, ADMM coordination, LightGBM imputer training and cross-validation, the communication-failure sweep, power-flow validation, and reporting — and each stage emits a governed JSON report following a fixed schema. Every report carries the same envelope fields (`report_id`, `schema_version`, `created_at`, `source_domain`, `inputs`, `artifacts`, `summary`, `validation`), and each recorded artifact stores its byte count and SHA-256 digest, so every number in this paper is traceable to a specific file and run. A regression baseline pins the headline aggregate peaks, and the run is deterministic: fixed random seeds and a synthetic weather day remove any external data dependency, so a reader can re-run the study from its configuration and obtain baseline-matching results. We claim re-runnable, citable reproducibility on this single configuration — not statistical generality across feeders or climates, which the limitations discuss. Table 3 lists the eight stages and their principal governed outputs.

VI. RESULTS

A. AGGREGATE COORDINATION

Figure 1 shows the cold-day aggregate load. The uncoordinated profile exhibits the classic double peak — a morning and an evening heating surge — whereas the ideal ADMM coordination flattens it to a nearly constant ≈ 147.7 kW. The $\varrho = 0.8$ trace, in which most agents are imputed, recovers much of the flattening but leaves residual peaks where imputation error accumulates.

TABLE 3. Workflow stages and principal governed outputs.

#	Stage	Principal output
1	Weather generation	Synthetic cold-day temperature
2	Baseline heating	Per-home $h_{i,t}$ profiles
3	Background load	Fixed non-deferrable aggregate B
4	ADMM coordination	Ideal $x_{i,t}$, iteration count
5	Imputer training	LightGBM model + CV metrics
6	Communication sweep	Per- ϱ imputed aggregates
7	Power-flow validation	$V_{\min}$, line loading, verdict
8	Reporting	Figures, KPI tables, JSON report


**FIGURE 1.** Aggregate cold-day load: uncoordinated double-peak versus the flat ADMM-ideal profile (≈ 147.7 kW) and the $\varrho = 0.8$ imputed case with residual peaks. Here ϱ is the non-responsive fraction.

Table 4 quantifies the aggregate key performance indicators. Ideal coordination cuts the peak from 169.1 to 147.7 kW, a 12.7% reduction, and drives the peak-to-average ratio (PAR) from 1.146 down to 1.001 — essentially a flat load. A PAR of 1.001 is the practical floor: it means the coordinated aggregate is constant to within a tenth of a percent across the entire day, so the flattening objective in (2) has been met almost exactly. Energy cost falls slightly, from 371.0 to 364.8 CAD — a 1.66% reduction — because the daily-energy constraint conserves total heating; the saving comes from shifting load out of the evening-peak TOU window (0.158 CAD/kWh) toward off-peak hours (0.078 CAD/kWh) rather than from using less energy. The modest size of the cost reduction is itself a useful result, and an honest one: the objective in (2) flattens the aggregate (it penalizes departure from a constant level), it does not chase the TOU price signal. A purely flat profile still places substantial load in the priced evening window, so the bill moves only as a side effect of peak-shaving. With energy conserved, demand response of this kind buys capacity relief (a lower peak) far more than it buys energy savings, and the case for it must rest on the network benefit rather than the bill.

1) Comfort preservation

The flattening is achieved without sacrificing thermal comfort. The mean absolute heating deviation of the ideal coordinated schedule from the thermostat baseline is 0.298 kW

TABLE 4. Aggregate KPIs by scenario.

Scenario	Peak (kW)	PAR	Cost (CAD)	RMSE (kW)
Uncoordinated	169.1	1.1458	371.0	—
Coordinated ideal	147.7	1.0007	364.8	—
Imputed $\varrho=0.0$	147.7	1.0007	364.8	0.000
Imputed $\varrho=0.2$	147.7	1.0009	364.9	0.170
Imputed $\varrho=0.4$	147.8	1.0012	364.9	0.242
Imputed $\varrho=0.6$	148.1	1.0035	364.9	0.252
Imputed $\varrho=0.8$	159.4	1.0802	366.2	0.234

per home-step — small against the ≈ 5 kW heating level the homes operate at. This is comfort preserved by construction: the projection (4) keeps every step inside the $\pm 50\%$ band of its baseline and conserves each home’s daily heating energy, so the coordinator can only re-time heat, never withhold it. The result is a feeder-level flat profile that is, at the level of each individual home, only a gentle nudge of when the thermostat draws power.

The imputed scenarios track the ideal closely up to $\varrho = 0.6$ and only degrade materially at $\varrho = 0.8$, where the peak rises to 159.4 kW and PAR to 1.080. The per-scenario imputation RMSE (last column) is non-monotone in ϱ : it grows from 0.170 to 0.252 kW as the imputed population grows, then dips to 0.234 kW at $\varrho = 0.8$. This is an averaging effect — with 24 of 30 agents imputed, individual errors partially cancel in the aggregate — and it explains why the aggregate-level damage at $\varrho = 0.8$ comes less from raw imputation error than from the loss of the responsive agents whose true, flexible schedules were doing most of the flattening work.

B. IMPUTATION ACCURACY AND ROBUSTNESS

The LightGBM imputer is a competent forecaster of an unseen home’s heating. Under 5-fold leave-agents-out cross-validation it achieves an RMSE of 0.294 kW and an MAE of 0.247 kW, against a heating standard deviation of 0.575 kW. At roughly 51% of the natural variability, the imputer is clearly better than predicting the mean and is accurate enough that the coordinated aggregate barely moves for small non-responsive fractions, as seen in the per-scenario RMSE column of Table 4.

Figure 2 plots peak and PAR against ϱ . Both stay essentially flat through $\varrho = 0.6$ and then turn sharply upward at $\varrho = 0.8$, the signature of a robustness threshold rather than a gradual decay.

A secondary finding concerns tractability rather than outcome quality. The ADMM iteration count grows steeply with ϱ : 8 iterations for the ideal case, then 14, 21, 150, and 280 for $\varrho = 0.2, 0.4, 0.6, 0.8$ respectively (all runs converged). Communication failure thus degrades coordination tractability — the effort to reach consensus — well before it degrades the coordinated outcome, an early-warning signal an operator could monitor.

The shape of this curve is informative. From $\varrho = 0.0$ to $\varrho = 0.4$ the iteration count roughly doubles per step (8, 14,




FIGURE 2. Aggregate peak and PAR versus the non-responsive fraction ρ . The horizontal axis labelled “rho” in the plotted image denotes this fraction ρ , not the ADMM penalty ρ_{ADMM} .

21), an almost linear growth that the coordinator absorbs easily. Between $\rho = 0.4$ and $\rho = 0.6$ it jumps by a factor of seven (21 to 150), and by $\rho = 0.8$ it reaches 280 — 35 times the ideal effort. The knee of this curve sits at the same place as the feasibility knee in the network results below, near $\rho = 0.6$. The two are linked: as more agents are imputed, the aggregate the coordinator sees diverges from the aggregate that physically forms, the sharing residual becomes harder to drive to zero, and convergence slows. Iteration count is therefore a cheap, purely software-side proxy for the harder-to-observe physical-feasibility margin, because both respond to the same underlying quantity — the mismatch injected by imputation.

Read as a sequence of effects, the $8 \rightarrow 280$ growth says that communication failure degrades the solver before it degrades the outcome. The coordinator must work thirty-five times harder at $\rho = 0.8$ than in the ideal case to reconcile an aggregate that no longer matches reality, yet the aggregate it eventually produces is still only mildly worse (a 159.4 kW peak against the ideal 147.7 kW) until the very last sweep point. Convergence effort is thus a leading indicator: it rises while outcomes are still acceptable, giving an operator a purely computational signal — visible before any meter or power-flow check — that the coordination problem is becoming ill-conditioned and a feasibility breach may follow.

C. NETWORK VALIDATION

The closed-loop contribution is summarized in Figure 3 and Table 5. The uncoordinated aggregate loads the worst trunk line to 115% — a thermal violation — and drops the worst minimum voltage to 0.905 p.u. Ideal coordination clears the thermal overload: maximum line loading falls to 99.7% (no violation) and worst minimum voltage rises to 0.918 p.u. This is the core result: flattening the aggregate is not merely a statistical improvement, it restores feeder feasibility.

The cleared state is resilient to communication failure up to $\rho = 0.6$, where the worst line rides the 100.0% edge but does not violate. At $\rho = 0.8$ the accumulated imputation error re-introduces a 108% overload — the feeder re-violates.




FIGURE 3. Worst minimum bus voltage (left axis) and worst maximum line loading (right axis) versus the non-responsive fraction ρ — the headline feasibility figure. Line loading clears below 100% through $\rho = 0.6$ and re-violates at $\rho = 0.8$. The image’s “rho” axis is the non-responsive fraction ρ .

TABLE 5. Network feasibility on IEEE-33 (worst of day).

Scenario	Peak (MW)	$V_{\min}$ (p.u.)	Loading (%)	Violation
Uncoordinated	4.60	0.905	115.0	yes
Coordinated ideal	4.02	0.918	99.7	no
Imputed $\rho=0.0$	4.02	0.918	99.7	no
Imputed $\rho=0.2$	4.02	0.918	99.7	no
Imputed $\rho=0.4$	4.02	0.918	99.7	no
Imputed $\rho=0.6$	4.03	0.918	100.0	no (edge)
Imputed $\rho=0.8$	4.34	0.911	108.0	yes

The story is therefore 115% $\rightarrow$ 99.7% $\rightarrow$ 108%: coordination clears the overload, the clearance survives substantial non-responsiveness, and only collapses once most agents must be imputed. This identifies a critical non-responsive fraction between 0.6 and 0.8 beyond which imputation can no longer protect the feeder.

a: An honest note on voltage.

We report voltage as a severity measure, not a compliance verdict. On the *case33bw* benchmark the worst minimum voltage never reaches the 0.95 p.u. ANSI limit in any scenario — the feeder is inherently stressed — so we do not claim voltage compliance. Coordination nonetheless reduces under-voltage depth (0.905 to 0.918 p.u.), which is a genuine improvement in feeder condition. The clean cleared/re-violated headline comes from line loading, where the 100% thermal limit gives an unambiguous binary verdict.

D. FORECAST UNCERTAINTY QUANTIFICATION

The network verdict above rests on a single realization per non-responsive fraction: one failing-agent subset and the raw point forecast. To test whether that point estimate understates risk, we quantify two sources of uncertainty. The first is aleatory — which agents fall silent, since the feeder impact depends on the placement and schedule of the missing population. The second is imputation error in the reconstructed heating. For each fraction ρ we draw 200 random failing subsets; for each subset the imputed agents’ heating is perturbed by Gaussian forecast residuals with standard deviation equal to the imputer’s cross-validated




FIGURE 4. Monte-Carlo forecast-uncertainty sweep (200 draws per fraction). Worst maximum line loading mean with its P5–P95 band versus the non-responsive fraction ρ (left axis, with the 100% thermal limit), and the probability of a line violation as bars (right axis). The image’s “rho” axis is the non-responsive fraction ρ .

TABLE 6. Monte-Carlo line loading and violation probability by fraction (200 draws each; mean [P5, P95]).

ρ	Loading (%) [P5, P95]	$P(\text{violation})$
0.0	99.7	0.00
0.2	99.7 [99.7, 99.7]	0.00
0.4	99.7 [99.7, 99.8]	0.00
0.6	100.4 [99.8, 102.0]	0.56
0.8	107.5 [106.0, 109.2]	1.00

RMSE (0.294 kW, Section above), sharing-ADMM is resolved, and the resulting daily peak is mapped to the worst-of-day minimum voltage and maximum line loading through a precomputed peak-to-network curve. The mapping is valid because, with fixed bus-share injection, the worst-of-day network metrics are monotone functions of the daily peak; the curve reproduces the deterministic operating points exactly — the 169.1 kW uncoordinated peak, for instance, maps to 0.905 p.u. and 115.0% loading.

Figure 4 and Table 6 report the result over 200 draws per fraction (mean with [P5, P95] band). The probability of a thermal violation rises from 0 for $\rho \leq 0.4$, through 0.56 at $\rho = 0.6$, to 1.00 at $\rho = 0.8$. The single-realization analysis reported $\rho = 0.6$ as marginally compliant (99.98% loading), but once the uncertainty in which agents fail and in the forecast is accounted for, the worst line exceeds its thermal limit in 56% of draws: the point estimate understated the risk. The [P5, P95] band is tight at low ρ — forecast errors average out across a large responsive population — and widens sharply once a majority of agents are imputed and that averaging no longer holds. This reframes the operator takeaway: the critical transition is not a single deterministic fraction but a probabilistic risk threshold centered near $\rho = 0.6$, where roughly one realization in two breaches the feeder.

E. DISCUSSION

Three observations tie the results together. First, the aggregate, tractability, and feasibility curves all share the same

knee near $\rho = 0.6$: the iteration count, the peak, and the worst line loading all stay benign through $\rho = 0.6$ and break together at $\rho = 0.8$. This suggests that a single observable quantity — the divergence between the imputed and the realized aggregate — governs all three, and that an operator monitoring ADMM convergence effort has an early proxy for impending feeder violation. Second, the network result is not a restatement of the aggregate result: a 12.7% peak cut is a statistical improvement, but clearing a 115% overload to 99.7% is an operational state change from infeasible to feasible, and the two need not coincide on a different feeder or load allocation. Third, the failure at $\rho = 0.8$ is graceful rather than catastrophic — 108% loading, not a non-converging power flow — which indicates that the imputer keeps the system in a recoverable regime even when most agents are silent. Together these support the central claim that coordination quality should be measured by feeder feasibility, and that imputation extends the range of communication failure over which that feasibility holds.

VII. CONCLUSION

We presented a network-validated pipeline that coordinates cold-climate electric heating with sharing-ADMM, reconstructs non-responsive agents with a LightGBM imputer, and — the closing contribution — validates the coordinated aggregate on the IEEE 33-bus feeder with AC power flow. The closed loop turns abstract flattening metrics into a feeder verdict: ideal coordination cuts the aggregate peak by 12.7% and clears a 115% thermal line overload down to 99.7%, while lifting the worst minimum voltage from 0.905 to 0.918 p.u. The cleared state is resilient to communication failure up to a non-responsive fraction of 0.6 and re-violates at 0.8 (108% loading). A Monte-Carlo analysis over which agents fail and the forecast error reframes this critical fraction as a probabilistic risk threshold rather than a single deterministic point: the probability of a thermal violation rises from 0 ($\rho \leq 0.4$) through 0.56 at $\rho = 0.6$ to 1.0 at $\rho = 0.8$, so the deterministic “edge” compliance at $\rho = 0.6$ understated the true risk. A secondary finding is that ADMM iteration count explodes with ρ (8 to 280) well before the outcome degrades, offering an early-warning signal.

The study has clear limitations. It covers a single cold day and a single feeder; the communication-failure model is a stochastic non-responsive fraction rather than a protocol-level simulation of delays, drops, and retransmissions; and the *case33bw* voltage profile is inherently stressed, so voltage is reported as a severity measure rather than a compliance test. Future work will extend the loop to multiple feeders and days, study the spatial placement of agents on the feeder (which determines how much each contributes to the binding constraint), and couple the coordinator to a flexibility market so that the cleared headroom can be priced and settled.

REFERENCES

- [1] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, "Distributed optimization and statistical learning via the alternating direction method of multipliers," *Foundations and Trends in Machine Learning*, vol. 3, no. 1, pp. 1–122, 2011.
- [2] T. Erseghe, "Distributed optimal power flow using ADMM," *IEEE Transactions on Power Systems*, vol. 29, no. 5, pp. 2370–2380, 2014.
- [3] S.-C. Tsai, Y.-H. Tseng, and T.-H. Chang, "Communication-efficient distributed demand response: A randomized ADMM approach," *IEEE Transactions on Smart Grid*, vol. 8, no. 3, pp. 1085–1095, 2017.
- [4] D. H. Nguyen, T. Narikiyo, and M. Kawanishi, "Optimal demand response and real-time pricing by a sequential distributed consensus-based ADMM approach," *IEEE Transactions on Smart Grid*, vol. 9, no. 5, pp. 4964–4974, 2018.
- [5] Y. Zheng, Y. Song, D. J. Hill, and Y. Zhang, "Multiagent system based microgrid energy management via asynchronous consensus ADMM," *IEEE Transactions on Energy Conversion*, vol. 33, no. 2, pp. 886–888, 2018.
- [6] F. Etedadi Aliabadi, S. Kelouwani, K. Agbossou, N. Henao, F. Laurencelle, and S. S. Hosseini, "Hierarchical transactive home energy management system groups coordination through multi-level consensus sharing-based distributed ADMM," *Sustainable Energy, Grids and Networks*, vol. 39, p. 101460, 2024.
- [7] S. Mhanna, G. Verbič, and A. C. Chapman, "Adaptive ADMM for distributed AC optimal power flow," *IEEE Transactions on Power Systems*, vol. 34, no. 3, pp. 2025–2035, 2019.
- [8] D. Gebbran, S. Mhanna, A. C. Chapman, and G. Verbič, "Multiperiod DER coordination using ADMM-based three-block distributed AC optimal power flow considering inverter volt-var control," *IEEE Transactions on Smart Grid*, vol. 14, no. 4, pp. 2874–2889, 2023.
- [9] R. Mieth and Y. Dvorkin, "Online learning for network constrained demand response pricing in distribution systems," *IEEE Transactions on Smart Grid*, vol. 11, no. 3, pp. 2563–2575, 2020.
- [10] F. Shen, Q. Wu, X. Jin, B. Zhou, C. Li, and Y. Xu, "ADMM-based market clearing and optimal flexibility bidding of distribution-level flexibility market for day-ahead congestion management of distribution networks," *International Journal of Electrical Power & Energy Systems*, vol. 123, p. 106266, 2020.
- [11] Z. Ma, N. Yang, S. Zou, and Y. Shao, "Charging coordination of plug-in electric vehicles in distribution networks with capacity constrained feeder lines," *IEEE Transactions on Control Systems Technology*, vol. 26, no. 5, pp. 1917–1924, 2018.
- [12] E. Dall'Anese, K. Baker, and T. Summers, "Chance-constrained AC optimal power flow for distribution systems with renewables," *IEEE Transactions on Power Systems*, vol. 32, no. 5, pp. 3427–3438, 2017.
- [13] A. Hassan, R. Mieth, M. Chertkov, D. Deka, and Y. Dvorkin, "Optimal load ensemble control in chance-constrained optimal power flow," *IEEE Transactions on Smart Grid*, vol. 10, no. 5, pp. 5186–5195, 2019.
- [14] P. N. Brown, H. P. Borowski, and J. R. Marden, "Are Multiagent Systems Resilient to Communication Failures?" *arXiv preprint arXiv:1710.08500*, 2017.
- [15] C. Gao, X. He, H. Dong, H. Liu, and G. Lyu, "A survey on fault-tolerant consensus control of multi-agent systems: trends, methodologies and prospects," *International Journal of Systems Science*, vol. 53, no. 13, pp. 2800–2813, 2022.
- [16] J. Xu, H. Sun, and C. J. Dent, "ADMM-based distributed OPF problem meets stochastic communication delay," *IEEE Transactions on Smart Grid*, vol. 10, no. 5, pp. 5046–5056, 2019.
- [17] J. Wu, T. Yang, D. Wu, K. Kalsi, and K. H. Johansson, "Distributed optimal dispatch of distributed energy resources over lossy communication networks," *IEEE Transactions on Smart Grid*, vol. 8, no. 6, pp. 3125–3137, 2017.
- [18] N. Bastianello, R. Carli, L. Schenato, and M. Todescato, "Asynchronous distributed optimization over lossy networks via relaxed ADMM: Stability and linear convergence," *IEEE Transactions on Automatic Control*, vol. 66, no. 6, pp. 2620–2635, 2021.
- [19] A. Gusrialdi and Z. Qu, "Resilient distributed optimization against cyber-attacks," *IEEE Control Systems Letters*, vol. 7, pp. 3956–3961, 2023.
- [20] W. Lee and B. Park, "Examination of learning-augmented approaches for initializing distributed optimal power flow with consensus ADMM," *IEEE Access*, vol. 14, pp. 7600–7615, 2026.
- [21] M. Buwaneswaran, M. U. Danish, and K. Grolinger, "Conditional denoising diffusion model for missing value imputation in smart meter data," *Engineering Applications of Artificial Intelligence*, vol. 181, p. 115403, 2026.
- [22] R. Razavi-Far, M. Farajzadeh-Zanjani, M. Saif, and S. Chakrabarti, "Correlation clustering imputation for diagnosing attacks and faults with missing power grid data," *IEEE Transactions on Smart Grid*, vol. 11, no. 2, pp. 1453–1464, 2020.
- [23] Y. Yuan, K. Dehghanpour, and Z. Wang, "Mitigating smart meter asynchrony error via multi-objective low rank matrix recovery," *IEEE Transactions on Smart Grid*, vol. 12, no. 5, pp. 4308–4317, 2021.
- [24] S. Ben Taieb and R. J. Hyndman, "A gradient boosting approach to the kaggle load forecasting competition," *International Journal of Forecasting*, vol. 30, no. 2, pp. 382–394, 2014.
- [25] J. R. Lloyd, "GEFCom2012 hierarchical load forecasting: Gradient boosting machines and gaussian processes," *International Journal of Forecasting*, vol. 30, no. 2, pp. 369–374, 2014.
- [26] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "Lightgbm: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [27] M. V. Gasca, F. Ibáñez, and D. Pozo, "Flexibility quantification of thermostatically controlled loads for demand response applications," *Electric Power Systems Research*, vol. 202, p. 107592, 2022.
- [28] M. Ali, J. Jokisalo, K. Siren, and M. Lehtonen, "Combining the demand response of direct electric space heating and partial thermal storage using LP optimization," *Electric Power Systems Research*, vol. 106, pp. 160–167, 2014.
- [29] B. Baeten, F. Rogiers, and L. Helsen, "Reduction of heat pump induced peak electricity use and required generation capacity through thermal energy storage and demand response," *Applied Energy*, vol. 195, pp. 184–195, 2017.
- [30] V. Lakshmanan, M. Marinelli, A. M. Kosek, P. B. Nørgård, and H. W. Bindner, "Impact of thermostatically controlled loads' demand response activation on aggregated power: A field experiment," *Energy*, vol. 94, pp. 705–714, 2016.
- [31] C. Guzman, K. Agbossou, and A. Cardenas, "Energetic macroscopic representation of an electrically heated building with electric thermal storage and heating control for peak shaving," *Energy and Power Engineering*, vol. 7, no. 4, pp. 144–153, 2015.
- [32] S. H. Nagarsheth, N. Henao, K. Agbossou, and M. Fournier, "Energy management set-up to exploit the flexibility potential of a multi-unit residential apartment building in a cold climate region," *IFAC-PapersOnLine*, vol. 58, no. 13, pp. 143–148, 2024.
- [33] D. Toquica, K. Agbossou, R. Malhamé, N. Henao, S. Kelouwani, and M. Fournier, "A recommender system for predictive control of heating systems in economic demand response programs," *IEEE Open Journal of Industry Applications*, vol. 3, pp. 79–89, 2022.
- [34] M. Kaheni, A. Pilloni, G. Serra Ruda, E. Usai, and M. Franceschelli, "Distributed asynchronous greedy control of large networks of thermostatically controlled loads for electric demand response," *IEEE Control Systems Letters*, vol. 7, pp. 169–174, 2023.
- [35] Z.-Q. Liu, Z.-W. Liu, L. Ye, D.-X. He, and G. Wen, "Distributed framework for dynamic coordination of thermostatically controlled loads: A resilient consensus approach," *Automatica*, vol. 189, p. 112974, 2026.
- [36] M. E. Baran and F. F. Wu, "Network reconfiguration in distribution systems for loss reduction and load balancing," *IEEE Transactions on Power Delivery*, vol. 4, no. 2, pp. 1401–1407, 1989.
- [37] L. Thurner, A. Scheidler, F. Schäfer, A. Probst, C. Spalthoff, M. Pock et al., "pandapower—an open-source python tool for convenient modeling, analysis, and optimization of electric power systems," *IEEE Transactions on Power Systems*, vol. 33, no. 6, pp. 6510–6521, 2018.